

HotRing: A Hotspot-Aware In-Memory Key-Value Store


Jiqiang Chen, Liang Chen, Sheng Wang, Guoyun Zhu,
Yuanyuan Sun, Huan Liu, and Feifei Li, Alibaba Group

FAST 2020

Background



➤ In-memory Key-Value Store

- Cache frequently accessed data for faster access and scalability
- The essential component in real world
- KV Store: ***Redis; Memcached***




Hotspot Issue

- A small fraction of data which are frequently accessed in a highly-skewed workload
 - Skewed workload: the same key is queried in the extreme short time
 - Daily cases: **50% of accesses for only 1% items**
 - Extreme cases: **90%**



Daily distributions

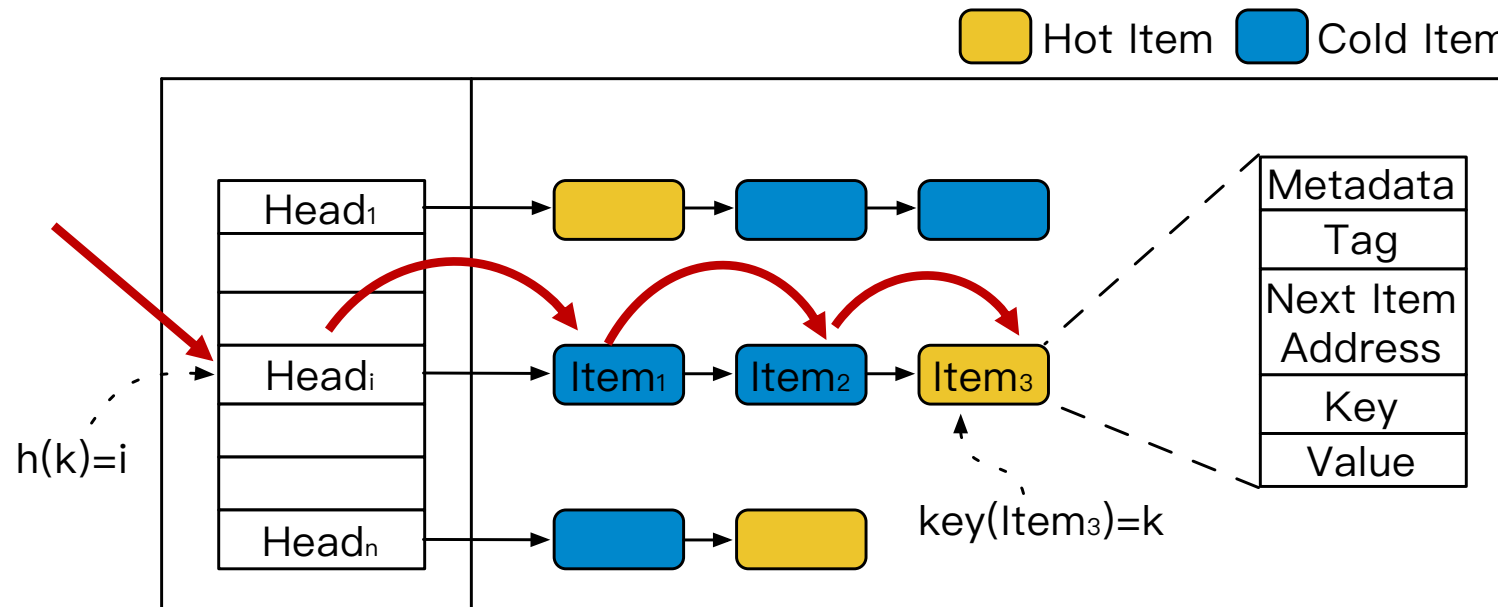


Extreme distributions

Typical Hash Index Structure

➤ Hash table with collision chain for each entry

- Item close to the tail requires more memory accesses (e.g., Item₃)
- Hotter item should be close to the head pointer for faster read



Cannot be aware of hotspot

Main Idea

- **How to reduce memory access overhead for hotspots?**
 - Make hotspot *closer to the head pointer*
- **How to identify the hotspot and its dynamic changes?**
 - Based on request frequency or access cost
- **How to provide efficient concurrent accesses at the same time?**
 - Using locks or non-locks schemes

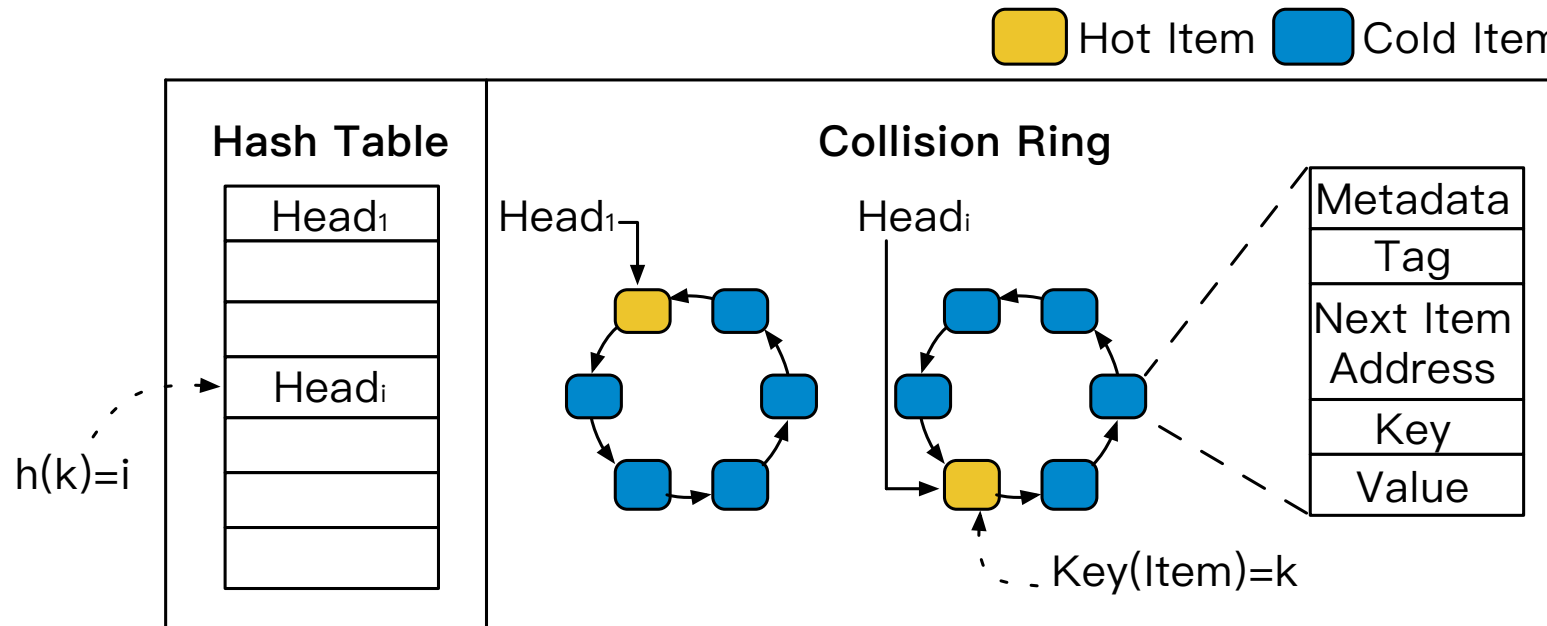
Contributions

- HotRing, an *ordered-ring hash structure* to enable hotspot-aware design
 - Move the head pointer closer to the hot item
 - Adopt a lightweight strategy to detect hotspot shifts at run-time
- Lock-free design for concurrency
 - Including hotspot shift detection, head pointer movement and ordered-ring rehash

HotRing

➤ The collision ordered-ring

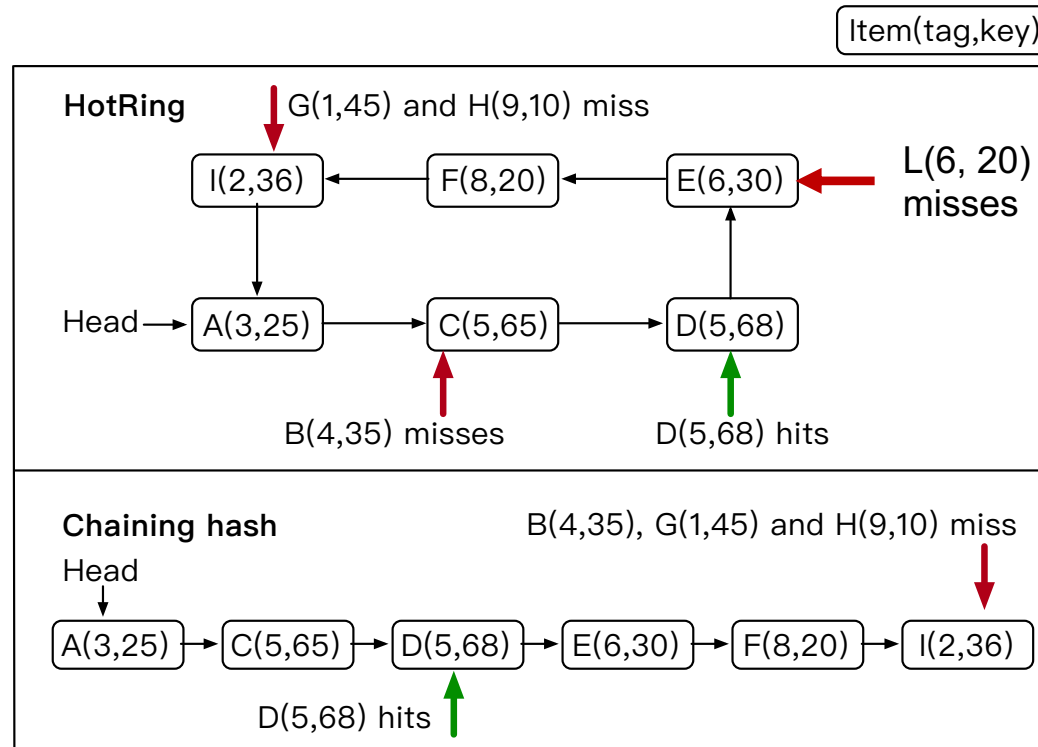

- Make the head pointer closer to the hot item
- Utilize ordered feature to faster lookup process (i.e., use tag and key)



HotRing

➤ The collision ordered-ring


- $\text{order}_k = (\text{tag}_k, \text{key}_k)$
- Use the order to faster lookup process or termination



Random Movement Strategy

➤ Based on random requests

- Adjust head pointer to the hotspot after every R requests (e.g., $R = 5$)
- Work well for a single hotspot
- Low accuracy for multiple hotspots but faster reaction




Reaction delay:
the time we detect the
hotspot after its
occurrence

Statistical Sampling Strategy

➤ Based on sampling


- Launch a new sampling after every R requests
- Record counts in the head pointer as well as in the item
- (The average memory access cost = $\text{count} * \text{distance}$)_{min}



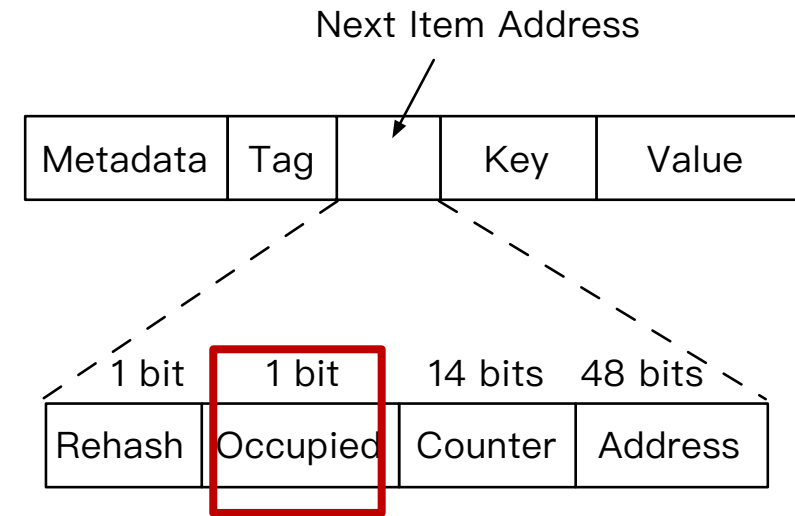
Lock-free Operations

➤ Adopt lock-free to support concurrency

- HotRing supports hotspot shift detection, head pointer movement and ordered-ring rehash
- Utilize occupied bit



Concurrency issue




Item format


Lock-free Rehash

➤ Rehash to increase data volume


- Access overhead is used to trigger rehashing instead of length for hotspot-awareness
- Tag range: $[0, T/2)$ and $[T/2, T)$, $T = 2^{(n-k)}$




Share highest bit of tag



Initialization



Split



Deletion

Experimental Setup

➤ Datasets:

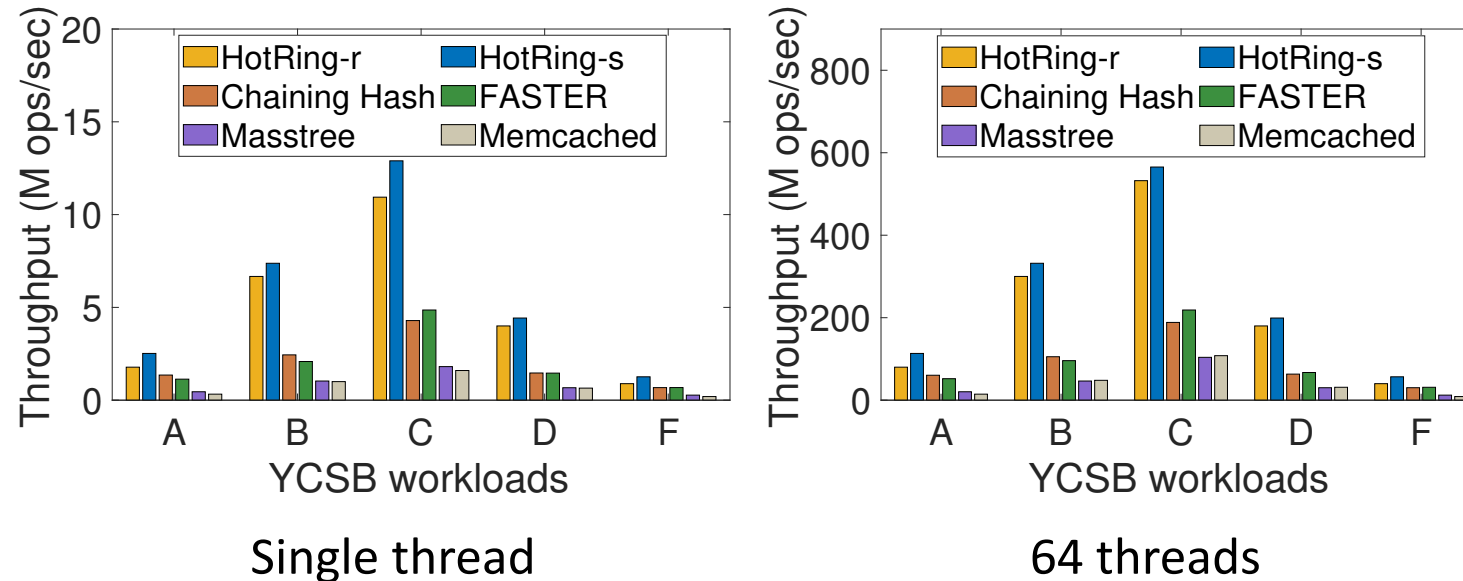
- YCSB workloads (Except for E)
- Key-value size: 8 bytes
- 250 millions of loaded keys

➤ Comparison:

- **Baseline:**
 - Lock-free chaining hash & FASTER
 - Masstree & Memcached
- **HotRing-r**
- **HotRing-s**

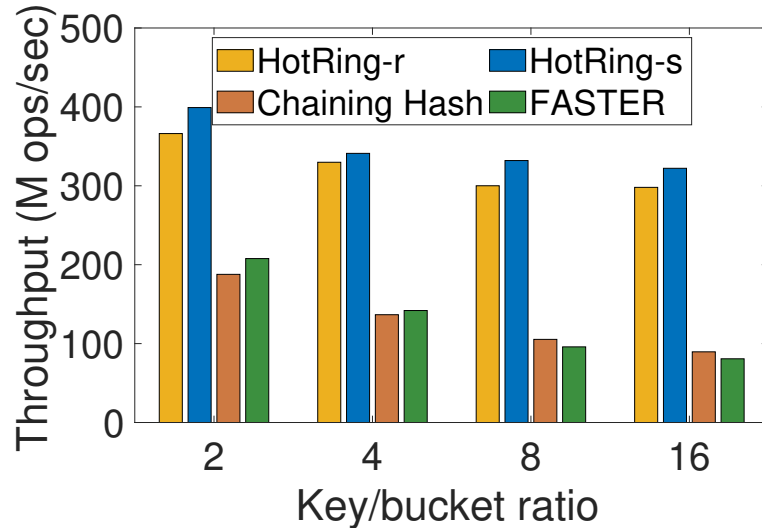
System	CentOS 7.4 OS with Linux 3.10 kernel
CPU	2.50GHz Intel Xeon(R) E5-2682 v4 * 2 (64 thread in total)
Cache Alignment	64B
Main Memory	32GB 2133MHz DDR4 DRAM * 8

Throughput




- HotRing outperforms other approaches by **2.10X-7.75X**
- Especially for B(95% read) and C(100% read)
 - HotRing-r is about **7% worse** than HotRing-s

Overall Performance



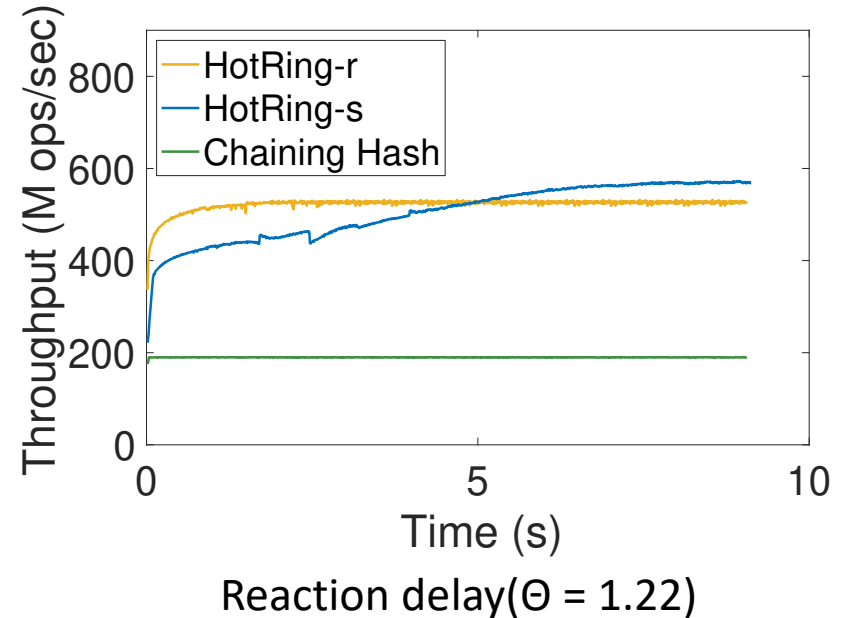

Impact of chain length



Hotspot awareness

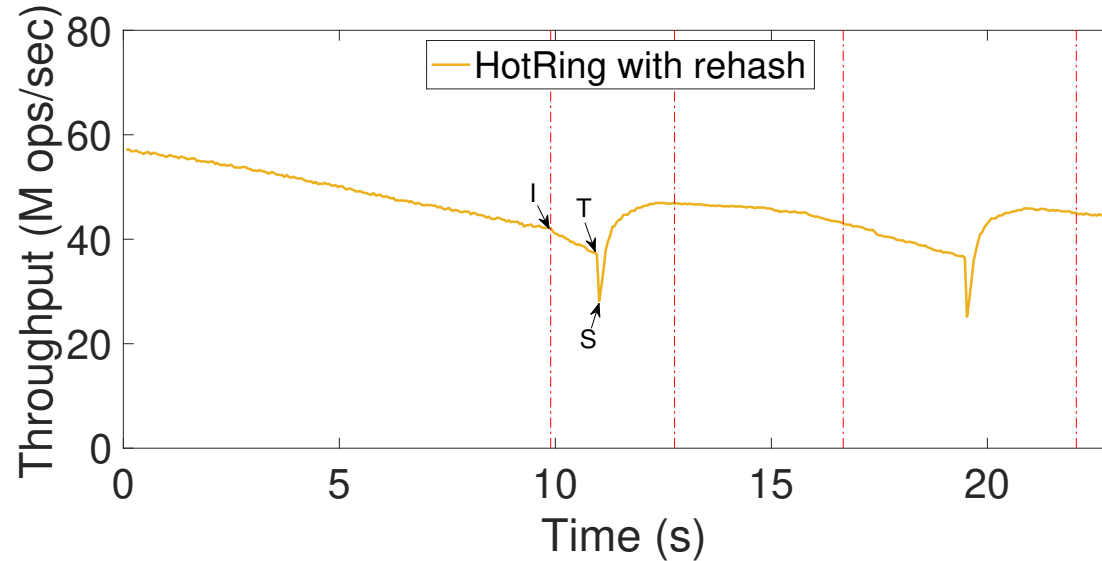
- HotRing's throughput increases from **1.3X to 3.91X** compared with FASTER since it puts hot items close to the head pointer (*less memory overhead*)
- HotRing improves up to 5X when the workload is more skewed

Micro-benchmark



- Non-HeadItems: HotRing-s has higher hotspot identification accuracy since more hotspot items are detected
- HotRing-r has faster reaction than HotRing-s (2 seconds to reach stable state)

Rehash Performance



Two consecutive rehash process ($\Theta = 1.22$)

- Split stage sets active flag to false, which blocks any requests for this collision ring.
- The short-term drops during rehash are because of the temporary lack of hotspot awareness when the new hash table starts to work.

Conclusion

➤ HotRing

- Hash Index with the ordered collision ring
- Utilize the pair of tag and key for ordering and fast lookup
- Solve the hotspot shift issue with hotspot-awareness strategy

➤ Concurrent support for HotRing

- Adopt lock-free method
- Support lock-free rehash operation