

# **Austere Flash Caching with Deduplication and Compression**

Qiuping Wang\*, Jinhong Li\*, Wen Xia# Erik Kruus^, Biplob Debnath^,  
Patrick P. C. Lee\*

\*The Chinese University of Hong Kong (CUHK)

#Harbin Institute of Technology, Shenzhen ^NEC Labs

**USENIX ATC'20**

# Background

- **Solid-state drives (SSDs) compared with hard disk drive (HDDs)**
  - Higher throughput; Lower latency; Better reliability
  - Performance, capacity and endurance are inversely proportional to price
- **Flash caching**
  - Accelerate HDD by caching frequently accessed blocks in flash
- **Deduplication & Compression**
  - Reduce storage & I/O overheads
  - Deduplication: In units of **chunks** (fixed-or-variable size)
  - Compression: In units of **bytes**

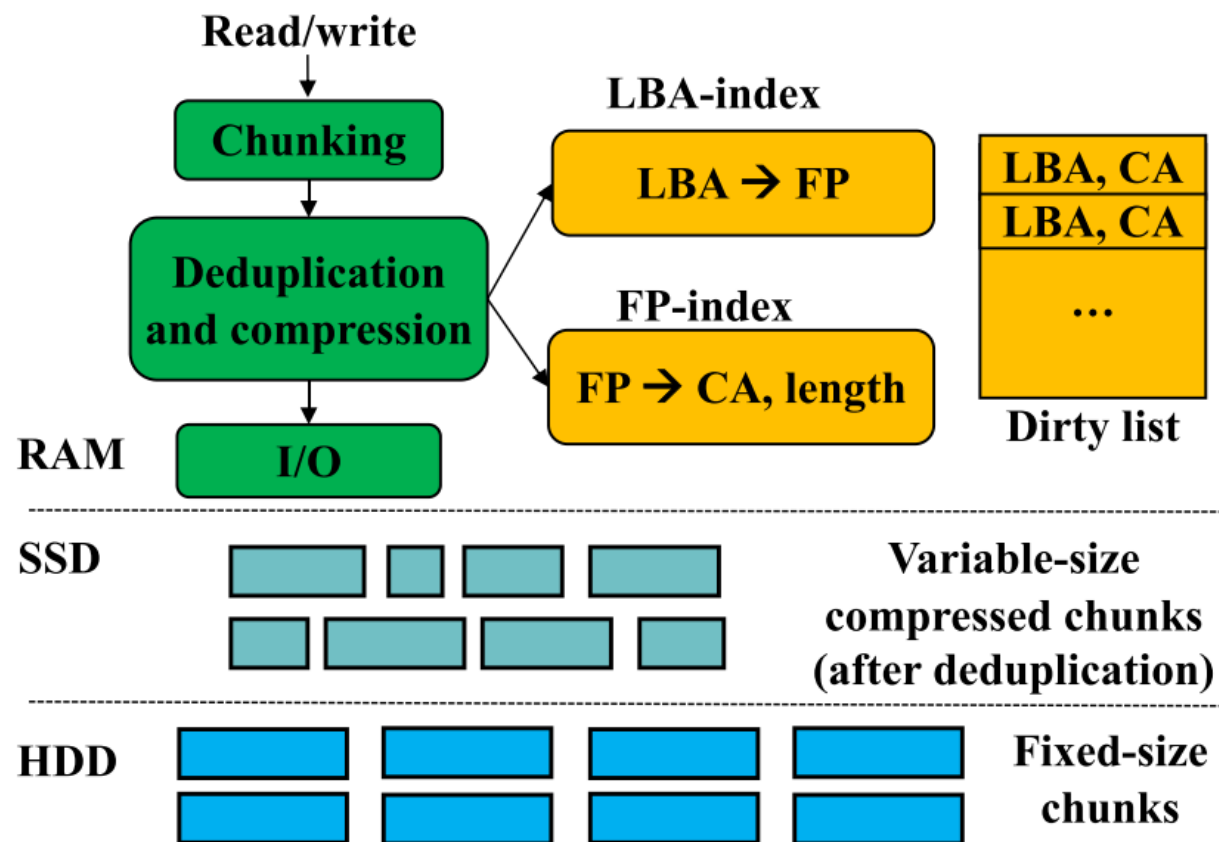
# Background & Problems

## ➤ Deduplicated and compressed flash cache

- **LBA**: chunk address in HDD
- **FP**: chunk fingerprint
- **CA**: chunk address in cache

## ➤ Memory amplification for index

- 32KB chunk, 512GB cache with 4TB HDD
- Conventional flash cache
  - LBA(8B)->CA(8B): 256MB
- With deduplication and compression
  - LBA(8B)->FP(20B): 3.5GB
  - FP(20B)->CA(8B)+Length(4B): 512MB
- At least **16x** memory amplification



# Contributions

## ➤ **AustereCache**: A flash cache support deduplication and compression with austere memory-efficient management

- Cache data structure based on **bucket** grouping and **prefix** index
  - Hash chunks to storage locations, no overhead for address mapping
  - Prefix index in the memory and the complete index in flash
- Fixed-size compressed data management
  - Compressing fixed-size chunks, **slice and pad** result to subchunks
  - No tracking for compressed length of chunks in memory
- Bucket-based cache replacement
  - Cache replacement per bucket
  - **Count-Min Sketch** for low-memory reference counting

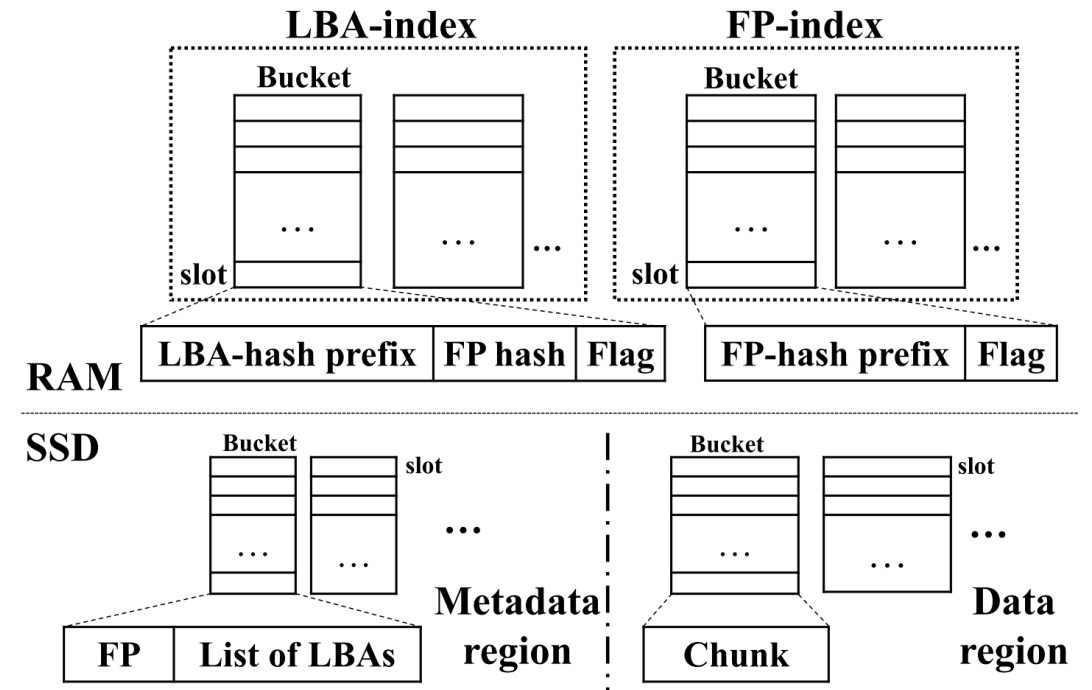
# System Architecture

## ➤ Data structure

- Use hashing to partition index and cache space
  - (RAM) **LBA-index** and **FP-index**
  - (SSD) **metadata region** and **data region**

## ➤ Layout

- Hash entries into equal-sized **buckets**
- Each bucket has fixed-number of **slots**

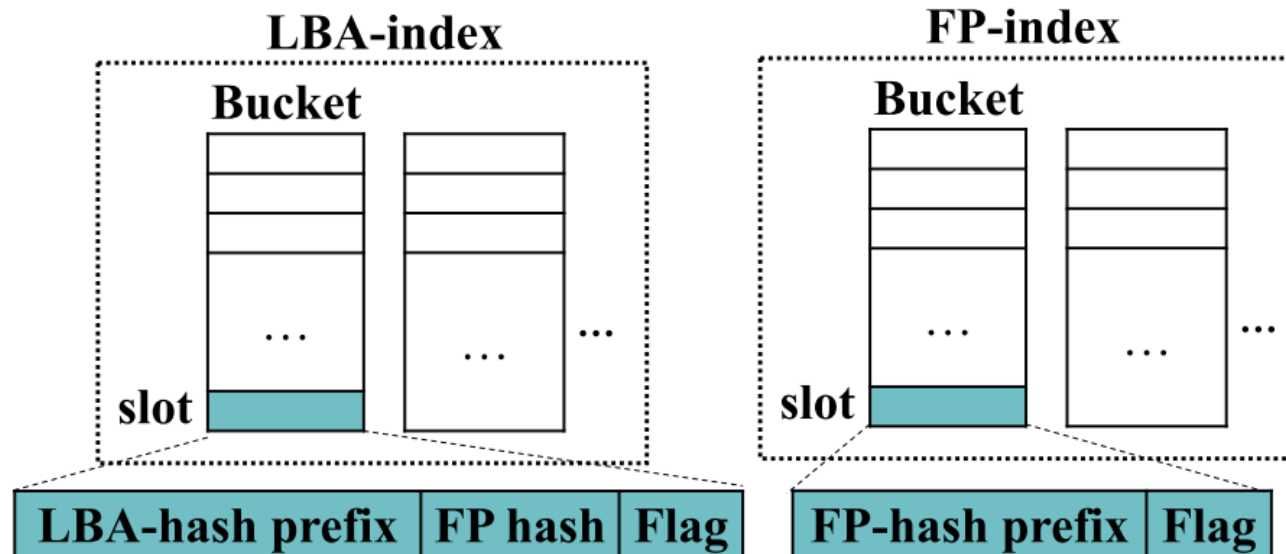


# LBA-index and FP-index

## ➤ Locate

- Buckets: hash suffixes
- Slot: hash prefixes

➤ Each slot in FP-index corresponds to a storage location in flash



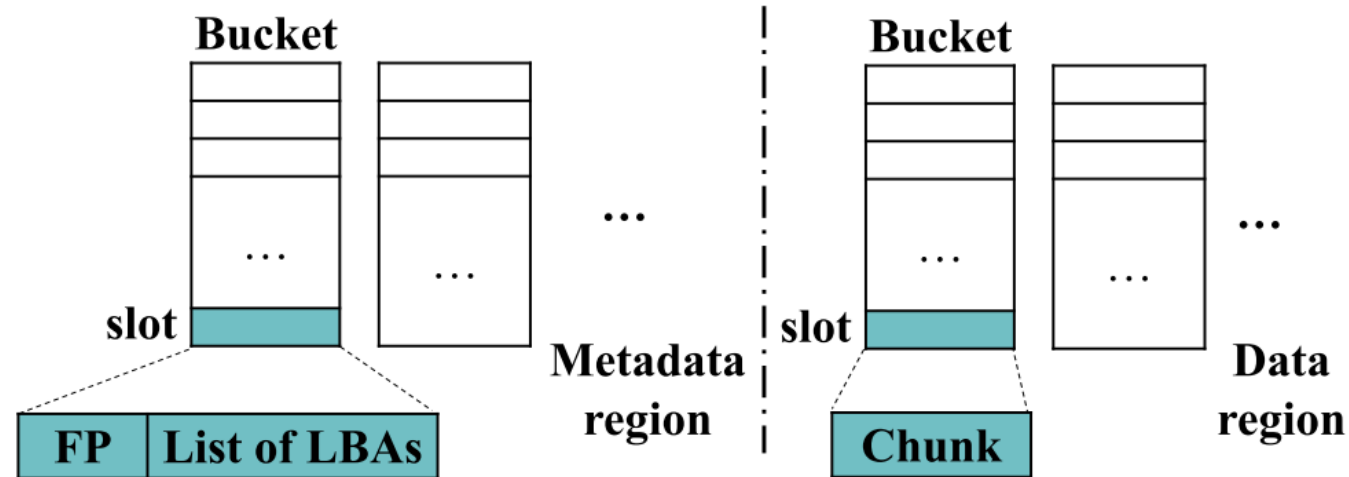
# Metadata and Data Regions

## ➤ Metadata region

- Store full FP and list of full LBAs (for validation against prefix collisions)

## ➤ Data region


- Locate by FP-index slot
- Store chunk content



# Fixed-size Compressed Data Management

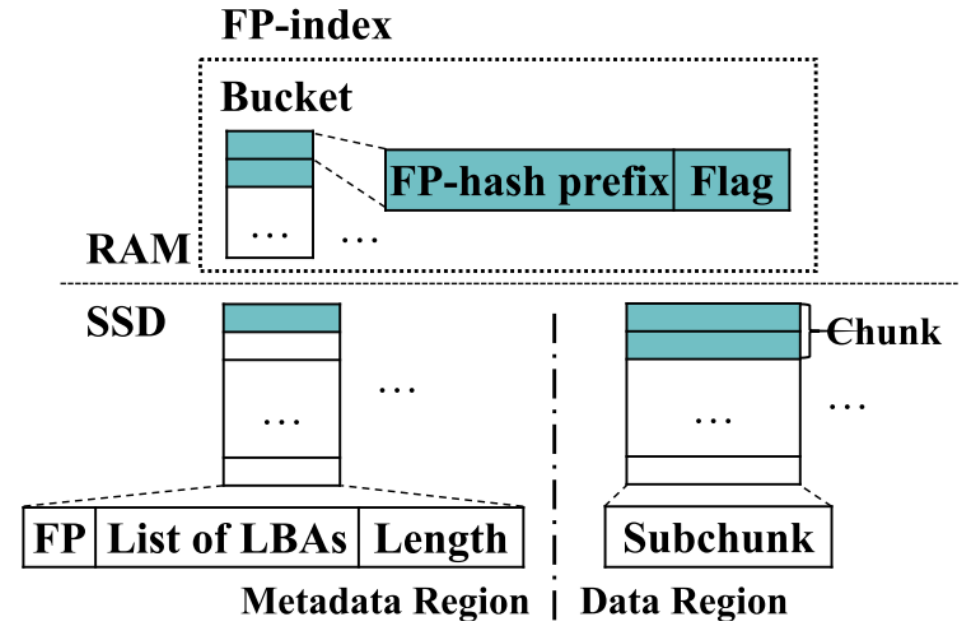
## ➤ Method

- Slicing and padding compressed chunk into fixed-size subchunks



## ➤ Layout

- One chunk use multiple consecutive slots
- Store chunk length in metadata region (no additional memory overhead)





# Fixed-size Compressed Data Management

## ➤ Advantages

- Compatible with bucketization (store each subchunk in one slot)
- Allow pre-chunk management for cache replacement

## ➤ Disadvantages

- The padding subchunk wastes storage space

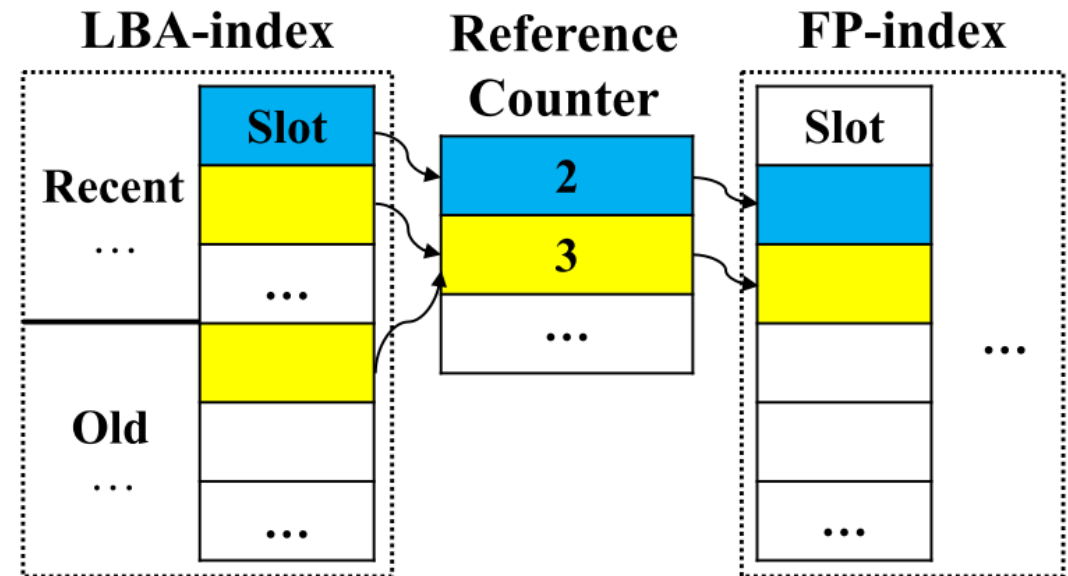
# Bucket-based Cache Replacement

➤ Cache replacement in each bucket independently


- LBA-index divided into recent/old

➤ Combine recency and deduplication

- LBA-index: least-recently-used policy
- FP-index: least-referenced policy
- Weighted reference counting based on recency in LBAs for FP-index

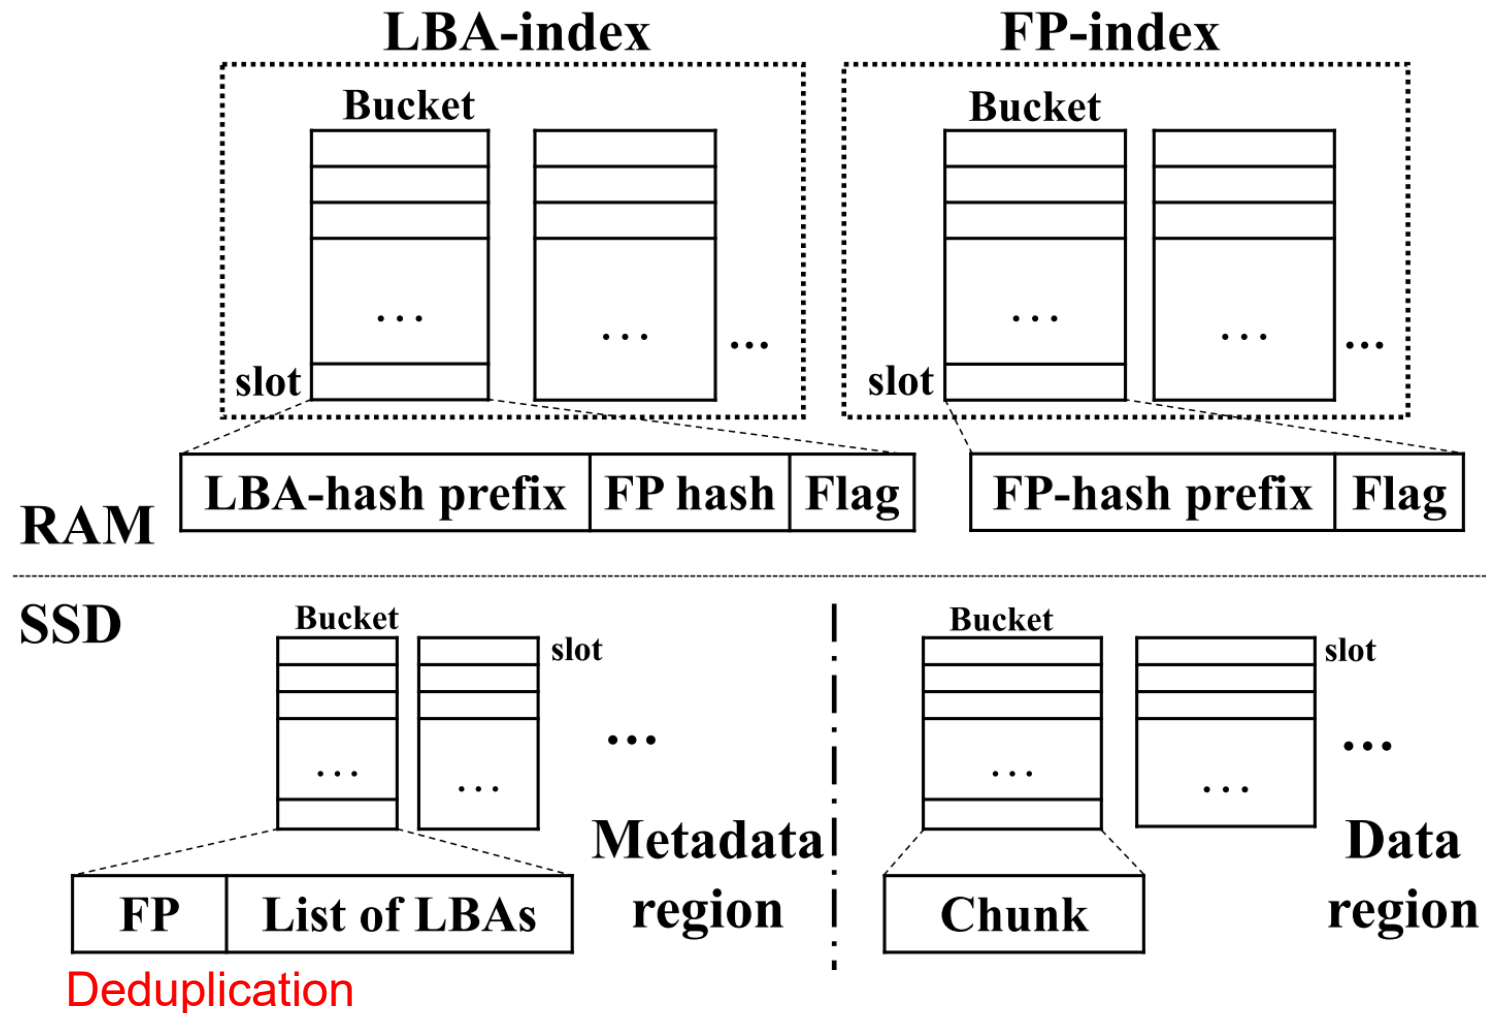


# Sketch-based Reference Counting



- High memory overhead for complete reference counting
  - One counter for every FP-hash
- Count-Min Sketch<sup>[Cormode 2005]</sup>
  - Fixed memory usage with provable error bounds

# Read, Write, and Deduplication Path



# Evaluation

## ➤ Traces


- FIU traces: WebVM, Homes, Mail
- Synthetic traces: varying I/O deduplication ratio and write-read ratio
  - I/O deduplication ratio: fraction of duplicate written chunks in all written chunks

## ➤ Schemes

- AustereCache:
  - AC-D Only deduplication
  - AC-DC deduplication & compression
- CacheDedup:
  - CD-LRU-D: LRU-based, only deduplication
  - CD-ARC-D: ARC-based, only deduplication
  - CD-ARC-DC: ARC-based, deduplication and compression

# Evaluation


## ➤ Memory overhead



- AC-DC incurs 87%~97% less memory usage than CD-ARC-DC
- AC-D incurs 70%~94% less memory usage than CD-LRU-D and CD-LRU-DC

# Evaluation


## ➤ Read hit ratios



- AC-D has up to 39.2% higher read hit ratio than CD-LRU-D, and similar read hit ratio as CD-ARC-D
- AC-DC has up to 30.7% higher read hit ratio than CD-ARC-DC

# Evaluation

## ➤ Write reduction ratios

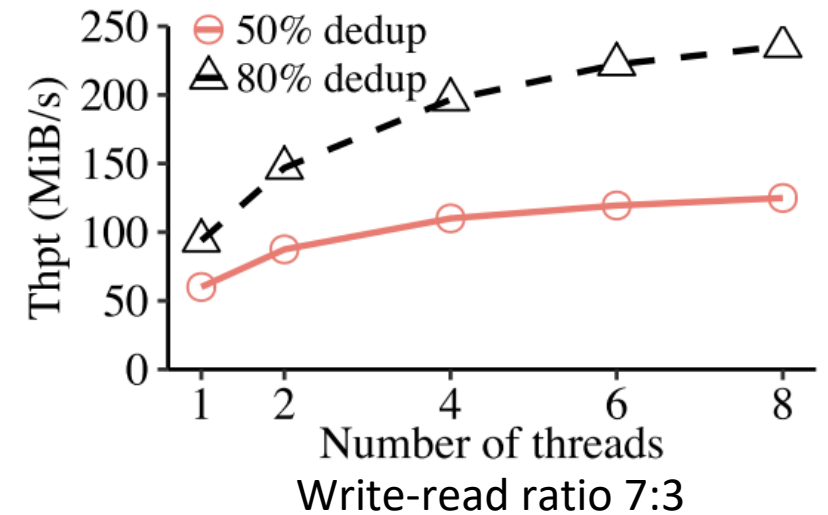



- AC-D is comparable as CD-LRU-D and CD-ARC-D
- AC-DC is slightly lower (by 7.7-14.5%) than CD-ARC-DC
  - Padding in compressed data management



# Evaluation

## ➤ Throughput



➤ AC-DC has highest throughput

➤ AC-D has slightly lower throughput than CD-ARC-D

- AC-D needs to access the metadata region during indexing

- 50% I/O dedup ratio: 2.08X
- 80% I/O dedup ratio: 2.51X

# Conclusion

- **AustereCache:** memory efficiency in deduplicated and compressed flash caching via:
- Bucketization
  - Fixed-size compressed data management
  - Bucket-based cache replacement

# Drawbacks

- The throughput is lower than HDD itself after adding the flash cache for single thread read or write.
- The size of the index structure is directly related to the hard disk's total size for caching.
- Use the slice and pad method to divide the compressed variable-length chunk into  $n$  fixed-length subchunks, which wastes part of the flash space.

# Q&A